

Network Programming

Lecture 1

LAN Programming – The Basics

Jonas Kunze

University of Mainz – NA62

Inverted CERN School of Computing, 24-25 February 2014

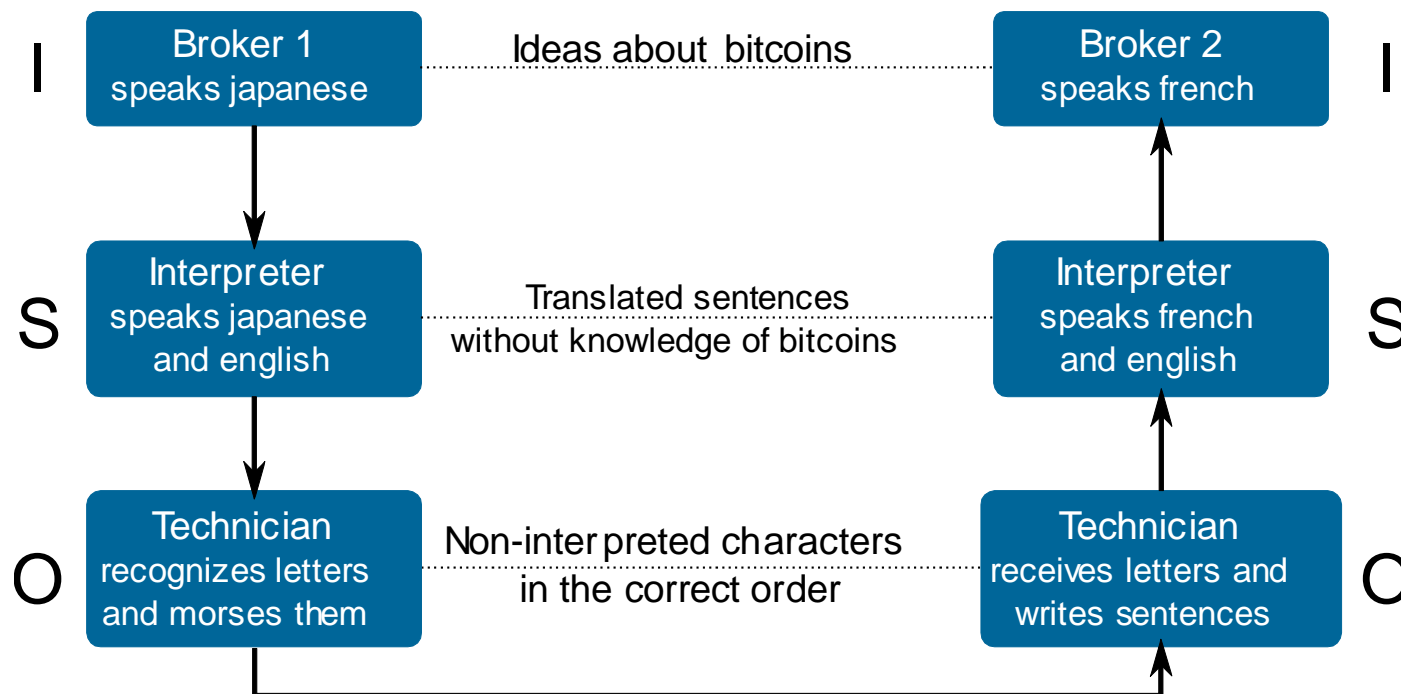
iCSC2014, Jonas Kunze, University of Mainz – NA62

Outline

- **Recap of the TCP/IP model**
 - ISO/OSI and TCP/IP
 - User Datagram Protocol (UDP)
 - Transmission Control Protocol (TCP)
- Network programming with BSD Sockets
 - Code snippets
 - Performance
- Alternatives to BSD Sockets
 - Network Protocols in User Space


The ISO/OSI reference model

- Communications protocols are divided into independent layers
- Every layer offers a service to the overlying layer



Interplay between OSI layers

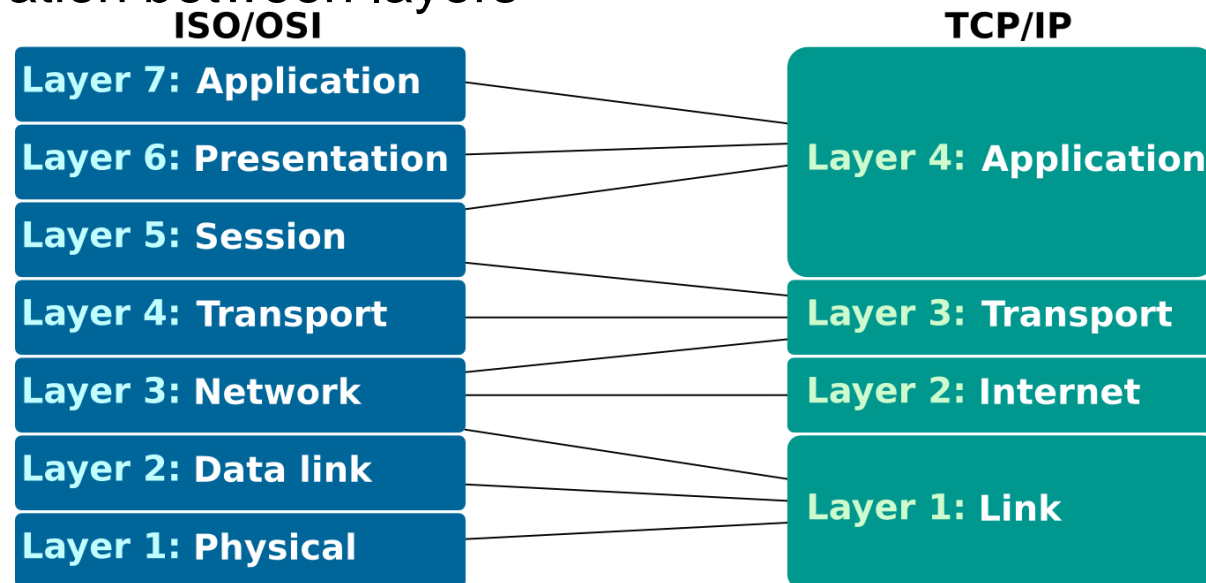
- **Every layer encapsulates the message into a Protocol Data Unit (PDU)**
 - PDUs typically consist of a Header and a Data section
- **Communication partners exchange PDUs by using the next lower layer**



- **Receiver unpacks PDUs in reverse order (like a stack)**


The TCP/IP model

- The ISO/OSI model is just a theoretical model with almost no implementation
- The most common communications protocols are part of the Internet Protocol Suite (TCP/IP)
 - Some ISO/OSI layers are merged
 - No strict separation between layers



User Datagram Protocol (UDP)

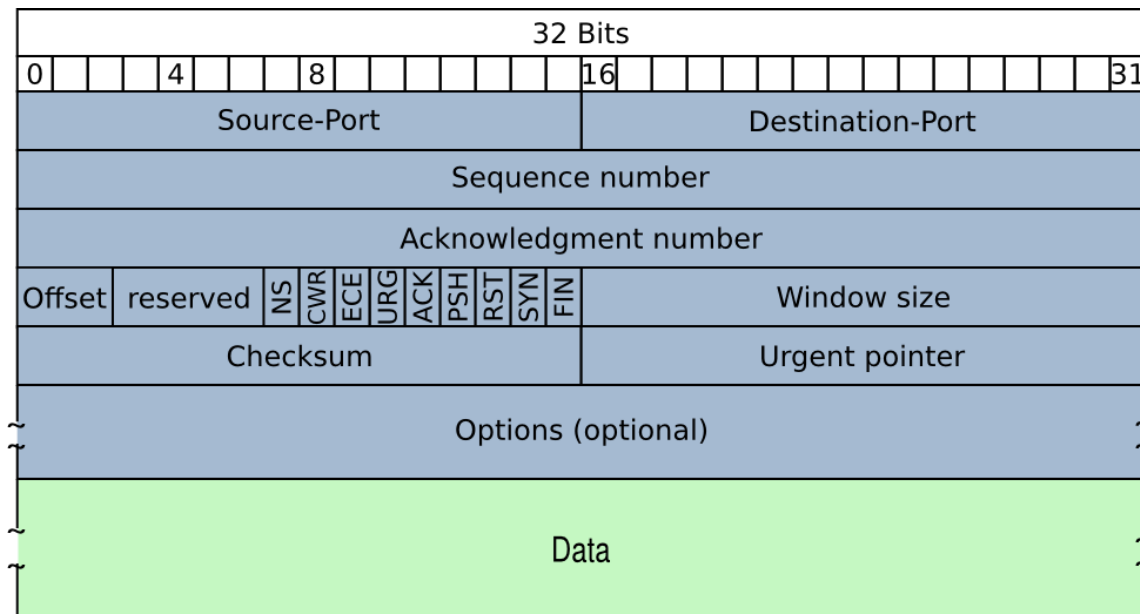
- UDP is connectionless and unreliable like IP



- Source-Port: The port of the process sending the datagram
- Destination-Port: The port number the datagram should be forwarded to
- Length: The length of the whole PDU in Bytes ($8 < \text{length} < 65535$)
- Checksum: Calculated with the whole PDU and data from the IP header

Transmission Control Protocol (TCP)

- Much more powerful and complex communication service than UDP
- Important application layer protocols based on TCP
 - World Wide Web (HTTP)
 - Email (SMTP)




Transmission Control Protocol (TCP)

- **TCP is reliable:**
 - Error-free: fragments are **retransmitted** in case they did not arrive at the destination (timeout)
 - **preserving order** without **duplicates**
- TCP is **connection oriented**
 - Connection establishment necessary before data can be sent
 - Connection defined by IP and **port number** (like UDP) of source and destination
 - Connections are always **point-to-point** and full-duplex
- It implements **flow control** and **congestion avoidance**
- Data is transmitted as an **unstructured byte stream**

TCP data flow

- A sends frame with SYN and random Sequence number X
- B acknowledges with $ACK=X+1$ and random Sequence number Y
- A acknowledges the reception
- A sends Z bytes
- B increases the sequence by Z to acknowledge the data reception
- Disconnection works like connection establishment but with FIN instead of SYN



Flow Control and Congestion Avoidance


- **Frames are only rarely dropped because of transmission errors (e.g. bit flip)**
 - Connections are typically either working without transmission errors or not at all
- Main reason for dropped frames are overloads of the **receiver** or the **network**

TCP implements two mechanisms to avoid overloading:

- **Flow control:** Avoids overloading of the receiver
- **Congestion avoidance:** Reduces the sending rate in case that fragments are dropped by the network

TCP's Flow Control: Sliding Window

- Each node has a receiving and sending buffer
- In each segment a node specifies how many bytes it can receive
 - Receiver window size: Number of free bytes in the receiving buffer
- If a node has sent as many **unacknowledged** bytes as the window size it will stop sending and wait for the next acknowledgment



- With each acknowledgment the window slides to the right


TCP's Congestion Avoidance

- **Congestion window:** Specifies the maximum number of bytes that may be sent without acknowledgment **depending on the network capacity**
- **Max bytes that may be sent = $\min(\text{sliding win}, \text{congestion win})$**

The congestion avoidance algorithm:

- Initialize the congestion window to typically 2 x MSS (slow start)
- Send until one of the two windows are filled
- If a segment is acknowledged: Increase the congestion window
 - Doubled until threshold reached, then linearly
- If acknowledgment timed out (frame dropped by network):
 - Set threshold to half the current congestion window and go back to slow start

TCP's Congestion Avoidance



Sending Buffer

- **When an application sends data chunks to the TCP stack two different approaches can be applied:**

1. Low latency

- Data chunks sent directly as they are
- Disadvantage: Many small IP packets will be transmitted (low efficiency)

2. High throughput

- Buffer data and send larger segments
- Higher latency but more efficient

Nagle's Algorithm

- **An algorithm to reach the high throughput approach:**
 - Send first chunk of data arriving at the TCP stack directly
 - Fill sending buffer with new incoming data without sending
 - If the buffer reaches the MSS : Send a new frame clearing the buffer
 - If all sent segments are acknowledged: Send a new frame clearing the buffer
- **Nagle's algorithm is used in almost all TCP implementations**
 - Can be deactivated to reduce latency (e.g. for X11 applications)

Switch off Nagle's Algorithm

- **This is only rarely necessary!**
- **Within your program:**

```
int flag = 0;
setsockopt(    socket,                /* socket affected */
               IPPROTO_TCP,          /* set option at TCP level */
               TCP_NODELAY,          /* name of option */
               (char ) &flag,        /* the actual value */
               sizeof(int));         /* length of option value */
```

- **System wide:**

```
echo 1 > /proc/sys/net/ipv4/tcp_low_latency
```


TCP vs UDP

- **TCP: A lot of bookkeeping and additional data transmission for acknowledgments**
- **UDP: Just sends the data as it is**

But...

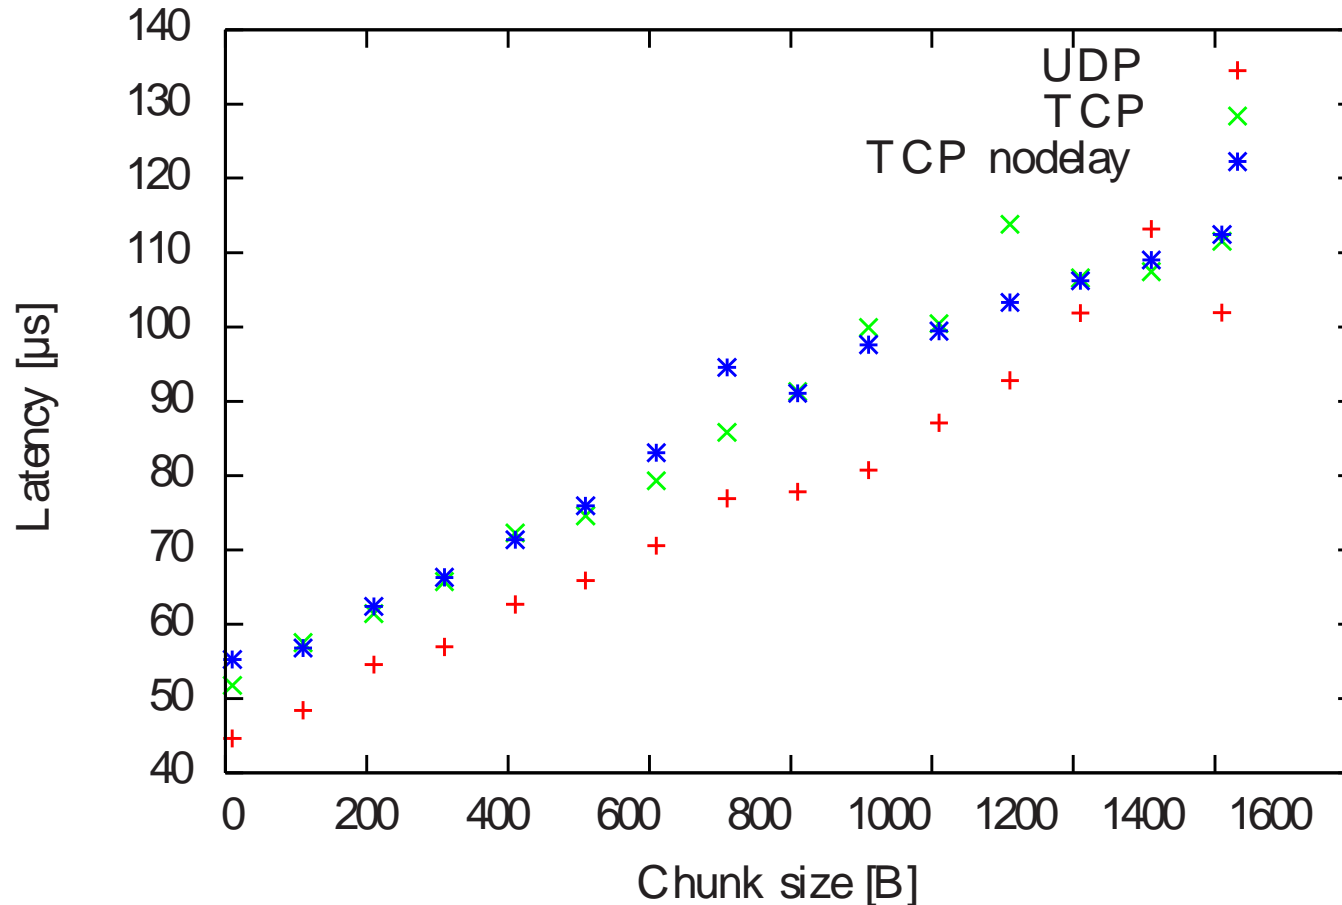
- **TCP: Flow control, congestion avoidance, Nagle's algorithm**

Typical rule of thumb:

- **TCP for high throughput, reliability and/or congestion avoidance**
- **UDP for low latency and broadcasts/multicasts (not possible with TCP)**

A Quick RTT Test

UDP vs TCP round trip times



This test was performed with hpcbench:
hpcbench.sourceforge.net

Outline

- Recap of the TCP/IP model
 - ISO/OSI and TCP/IP
 - User Datagram Protocol (UDP)
 - Transmission Control Protocol (TCP)
- **Network programming with BSD Sockets**
 - Code snippets
 - Performance
 - Interrupt Coalescing
 - NAPI
- Alternatives to BSD Sockets
 - Network Protocols in User Space

BSD Sockets

- **Linux supports TCP/IP as its native network transport**
- **BSD Sockets is a library with an interface to implement network communications using any TCP/IP layer below the application layer**
- **Important functions**
 - `socket()` opens a new socket
 - `bind()` assigns socket to an address
 - `listen()` prepares socket for incoming connections
 - `accept()` creates new socket for incoming connection
 - `connect()` connects to a remote socket
 - `send()` / `write()` sends data
 - `recv()` / `read()` receives data

TCP Code Snippet

Simple TCP socket accepting connections and receiving data:



```
socket = socket(AF_INET, SOCK_STREAM, 0);  
serv_addr.sin_family = AF_INET;  
serv_addr.sin_port = htons(8080);  
serv_addr.sin_addr.s_addr = INADDR_ANY;  
bind(socket, (struct sockaddr *) &serv_addr, sizeof(serv_addr));  
listen(socket, 5);  
connectionSocket = accept(socket, (struct sockaddr *) &cli_addr, &clilen);  
recv(connectionSocket, buffer, sizeof(buffer), 0);
```

Complete examples to be found at: <http://github.com/Jonaskunze>

TCP vs UDP: Throughput

Single threaded blocking sender and receiver, reliable network


- Small frames induce high CPU load → packet loss
- TCP achieves higher throughput



Down to the Kernel

- **When data arrives at the NIC:**
 - Data **copied** to kernel space (DMA)
 - NIC sends **interrupt**
 - Kernel **copies** data to the corresponding user space buffer (socket)
 - Kernel informs user space application

Linux Kernel Sockets



Interrupt Coalescing


- **Technique to reduce interrupt load**
- **Interrupts are held back until...**
 - ... a certain number of frames have been received...
 - ... or a timer times out
- **Now the kernel can process several frames at once**
 - Higher efficiency with just little increase of latency

```
# print current settings  
ethtool -c eth0
```

```
# change settings  
ethtool -C eth0 rx-usecs 0 # 0 is adaptive mode for many drivers  
ethtool -C eth0 rx-frames 12
```


Interrupt Coalescing

- Small values overload the CPU → Packet loss
- High values lead to buffer overflow → Packet loss



First bin shows adaptive mode

NAPI

- **An alternative to interrupts is polling:**
 - Kernel periodically checks for new data in the NIC buffer
 - High polling frequencies induce high memory loads
 - Low polling frequencies lead to high latencies and packet loss
- **NAPI: Linux uses both**
 - Interrupts per default
 - Polling in case of high data rates incoming


The kernel still needs to copy incoming data!

Outline

- Recap of the TCP/IP model
 - ISO/OSI and TCP/IP
 - User Datagram Protocol (UDP)
 - Transmission Control Protocol (TCP)
- Network programming with BSD Sockets
 - Code snippets
 - Performance
- **Alternatives to BSD Sockets**
 - Network Protocols in User Space
 - Example: pf_ring DNA
 - Reliability on top of UDP?
 - Reliability without acknowledgment

Network Protocols in User Space

- **Following approach can be implemented in the user space to avoid double copies**
 - NIC copies incoming data to a user space buffer (DMA)
 - The user space application polls the buffer
 - The user space application may enable interrupts for low data rates
 - The kernel is only used for the initialization
- **0% CPU used for accessing the data**



Example: pf_ring DNA

- **Proprietary user space driver by ntop**
- **Does not implement any protocol**
 - You need to implement them: ETH, IP, UDP, TCP, ARP, IGMP...
- **Compatible with all 1 GbE and 10 GbE NICs running on PCI-E**
- **Full line rate (1-10 GbE) with any frame size**
- **Round trip time below 5 μ s**
- **Hardware filtering (only Intel and Silicom NICs)**
 - Very efficient Intrusion prevention systems possible (Snort)
- **Other userspace drivers: Netmap, Intel DPDK, OpenOnload**

Reliability on top of UDP?

- **At CERN experiments most data senders are FPGAs**
 - Very fast in parallel jobs
 - Typically fully loaded by algorithms
 - Sometimes there's no space left for a fully implemented TCP/IP stack
- **I've seen many groups implementing reliable protocols on top of IP**
 - In most cases the result was TCP without flow and congestion control
- **Being compatible with TCP/UDP relieves the software developers**
 - You don't need to implement the protocol on the receiver side
 - Instead you can use standard libraries

Reliability without acknowledgment

- **Sometimes it's not even possible to store data until the acknowledgment is received**
 - You should use pure UDP in this case
- **As soon as datagrams are sent out you have to trust the network**
 - Make sure that you don't overload switches/routers/receiver nodes
 - Check every node whether frames are dropped

Switch/Router:

show interfaces ...

Linux:

cat /proc/net/udp

Summary

- **TCP is more than just reliable**
 - It implements a maximum efficient data transmission
- **BSD sockets provide a nice API for simple network programming**
 - For more complex architectures networking libraries are recommended
- **Linux' network sockets are not as efficient as they could be**
 - High performance network drivers provide efficient alternatives to BSD sockets but they generate additional work for the developer team